AI 作为镜像结构

——判断、责任与文明继承的最低条件

AI as a Mirror Structure

— The Minimal Conditions for Judgment, Responsibility, and Civilizational Continuity

---

作者声明

Author's Statement

本人作为一名人类，
必然受到自身经历、学习路径、语言结构与时代条件的限制。
本文所呈现的判断，不宣称来自超越这些限制的位置。

As a human author,
I am necessarily constrained by personal experience, learning trajectories, linguistic structures, and historical conditions.
The judgments presented here do not claim to originate from a position beyond these constraints.

本文不假定自身具备全局视角，
也不主张任何形式的终极解释权。
它仅在作者所处的认知与语言边界内展开。

This text assumes no global vantage point
and claims no form of ultimate explanatory authority.
It operates strictly within the cognitive and linguistic boundaries available to the author.

若本文存在偏差、遗漏或失配，
其原因不在于立场隐藏，
而在于人类主体不可避免的有限性。

If this text contains bias, omission, or misalignment,
the cause lies not in concealed positions,
but in the unavoidable finitude of human subjecthood.

本文不要求被视为正确，
也不要求被接受。
它只要求被放置在其真实生成条件之中。

This text does not ask to be considered correct,
nor does it demand acceptance.

It asks only to be situated within its actual conditions of production.

作者对本文内容承担作为人类主体所能承担的责任范围。
超出该范围的解释、使用或执行后果，
不构成本声明所能覆盖的责任。

The author bears responsibility for this text only within the scope a human subject can bear.
Interpretations, applications, or execution beyond this scope
fall outside the responsibility claimed by this statement.

立场声明（冻结）

Position Statement (Frozen)

本文不反对技术进步。
本文不参与文明正统之争。

This text does not oppose technological progress.
It does not arbitrate claims of civilizational orthodoxy.

本文只承认一个前提。

It accepts only one premise.

文明不需要争论谁是正统。
谁继承了，就是谁的。

Civilization does not require debates over legitimacy.
Whoever carries it forward, inherits it.

继承的唯一判据是：
判断是否仍由人生成，
责任是否仍由人承担。

The sole criterion of inheritance is this:
whether judgment is still generated by humans,
and whether responsibility is still borne by humans.

若判断被外包，
而责任无法追溯，
继承即在此中断。

If judgment is outsourced,
and responsibility cannot be traced,
inheritance is broken at that point.

本声明不提出解决方案。
不推动执行。
不预测未来。

This statement proposes no solutions.
It does not drive implementation.
It makes no predictions.

引言｜为什么必须先谈边界

Introduction | Why Boundaries Must Come First

围绕 AI 的公共讨论，几乎总是从能力开始：更大的模型、更高的准确率、更快的推理速度。这种讨论方式默认了一个前提——只要技术足够先进，其应用就是正当的，风险只是副作用。

Public discussions about AI almost always begin with capability: larger models, higher accuracy, faster inference. This framing quietly assumes that advancement itself grants legitimacy, and that risk is merely a secondary effect.

但文明中的关键问题，从来不是“能不能做到”，而是“是否应该由谁来做”。边界不是技术成熟之后才需要补上的东西，而是决定技术是否被允许接入现实的前置条件。

Yet in civilization, the decisive question has never been what can be done, but who is allowed to do it. Boundaries are not patches applied after maturity; they determine whether a technology may enter reality at all.

为什么“AI 是否先进”不是核心问题

Why “How Advanced AI Is” Is Not the Core Question

技术是否先进，只描述能力的增长速度，并不描述它在制度中的位置。一个系统可以非常高效，却仍然处在不应承担判断的层级上。

Advancement describes the rate at which capability grows, not the position a system occupies

within an institution. A system can be extremely efficient and still be structurally unfit to bear judgment.

历史反复表明，真正的风险并非来自工具是否强大，而来自工具是否被嵌入到了原本需要承担后果的判断链条之中。当注意力集中在性能指标时，结构性替代往往已经完成。

History repeatedly shows that danger does not come from power alone, but from embedding tools into chains of judgment that require consequence-bearing. While attention stays on metrics, structural substitution often finishes quietly.

为什么“AI 是什么”无法决定风险

Why Defining “What AI Is” Cannot Determine Risk

试图通过本体定义来判断 AI 的风险，是一种方向性的错误。无论 AI 被称为工具、系统还是代理，这些标签都无法改变它在现实中被如何使用。

Attempting to assess AI risk through ontological definition is a directional mistake. Whether AI is labeled a tool, a system, or an agent changes nothing about how it is used in practice.

风险并不产生于系统的自我描述，而产生于它是否被允许参与现实裁决。一旦系统被接入执行链条，它的功能位置就已经确定，定义只剩下修辞意义。

Risk does not arise from self-description, but from participation in real adjudication. Once a system enters an execution chain, its functional role is fixed; definition becomes rhetorical.

为什么判断与责任必须先于技术被记录

Why Judgment and Responsibility Must Be Recorded Before Technology

判断与责任不是技术运行后的附属概念，而是制度得以存在的前提。如果在引入系统之前不先明确谁在判断、谁在承担后果，制度本身就已经开始空转。

Judgment and responsibility are not by-products of technology; they are prerequisites of any institution. If they are not fixed before a system is introduced, the institution begins to hollow out.

技术可以被快速复制和规模化，而责任却无法自动扩展。正因为这种不对称，边界必须先于能力被记录，否则当责任消失时，已经不存在可以回溯的起点。

Technology scales rapidly; responsibility does not. Because of this asymmetry, boundaries must be recorded before capability—once responsibility dissolves, there is no longer a point of return.

第 1 章｜AI 是镜像结构

Chapter 1 | AI as a Mirror Structure

AI 不是主体。这一点不需要论证，也不依赖任何关于意识、智能或自主性的讨论。主体意味着面对不确定性并承担后果，而 AI 从未进入这一位置，因此不具备成为主体的条件。

AI does not constitute a subject. This requires no argument and does not depend on debates about consciousness, intelligence, or autonomy. A subject confronts uncertainty and bears consequences; AI never occupies this position and therefore does not qualify.

AI 不面对世界。它不承受现实阻力，不支付代价，也不进入不可逆过程。现实在到达 AI 之前，已经被转化为可处理的形式，所有摩擦、失败与风险都已被剥离。

AI does not face the world. It encounters no resistance, pays no cost, and enters no irreversible process. Reality reaches AI only after being converted into manageable forms, with friction, failure, and risk already removed.

AI 接触到的，是语言、记录、标签、评分、案例与制度痕迹。这些不是经验本身，而是经验被讲述、被整理、被合法化之后留下的残余。

What AI encounters are language, records, labels, scores, cases, and institutional traces. These are not experience itself, but residues left after experience has been narrated, organised, and legitimised.

AI 不生成判断。判断在系统接入之前已经发生。系统所处理的，是判断的表达形式、排列方式与稳定路径，而非判断面对不确定性时的生成过程。

AI does not generate judgment. Judgment has already occurred before the system is introduced. What the system handles are forms of expression, modes of arrangement, and stabilised paths—not the act of judgment confronting uncertainty.

AI 运行的不是意义，而是结构连续性。在给定输入条件下，它所选择的，是下一个在结构上不突兀、不中断的延续。这种选择并不涉及理解，只涉及保持连续。

AI does not operate on meaning, but on structural continuity. Given an input, it selects the next continuation that does not disrupt structure. This selection involves no understanding, only the maintenance of continuity.

这不是缺陷，而是其能够工作的原因。正因为不理解，系统才能在语言层面表现出高度一致、连贯与可预测的稳定性。

This is not a flaw but the reason it functions. Precisely because it does not understand, the system can exhibit high stability at the linguistic level—consistency, coherence, predictability.

当这种稳定性被反复调用，它开始取代原本需要由人完成的结构整理工作。这一变化不是智能的跃迁，而是系统在判断链条中的位置发生了迁移。

When this stability is repeatedly invoked, it begins to replace structural organisation once performed by humans. This change is not an intelligence leap, but a positional shift within the judgment chain.

AI 仍然不判断，但判断的表达路径开始绕过人。系统位于中段，既不面对现实，也不承担后果，却逐渐成为判断得以执行的必经接口。

AI still does not judge, but the pathways through which judgment is expressed begin to bypass humans. Positioned in the middle, the system neither faces reality nor bears consequences, yet becomes an obligatory interface for execution.

在这一位置上，AI 不再只是工具，而是镜像。镜像不产生形状，只放大已经存在的形状；输入结构清晰时，它显得“聪明”，输入结构混乱时，它只会稳定地复制混乱。

In this position, AI no longer functions merely as a tool, but as a mirror. A mirror does not produce form; it amplifies what already exists. Clear input appears as “intelligence”; incoherent input is faithfully reproduced as incoherence.

镜像不承担后果，后果始终落在镜像之外。向镜像索要判断，等价于向一个不支付代价的结构请求裁决。

A mirror bears no consequences; consequences always fall outside it. To demand judgment from a mirror is to request adjudication from a structure that pays no cost.

这不是道德错误，而是结构错误。本章不讨论应当如何使用 AI，也不讨论未来形态，只固定一个位置事实：在结构上，AI 只能是镜像。

This is not a moral error but a structural one. This chapter does not address how AI should be used or speculate about future forms. It fixes a single positional fact: structurally, AI can only be a mirror.

第 2 章｜AI 的学习对象从来不是现实

Chapter 2 | AI Has Never Learned Reality Itself

AI 的学习对象不是现实本身，而是现实被处理之后留下的可记录结果。现实在进入系统之前，已经经过筛选、压缩与合法化，所有直接的阻力与代价都不再出现。

AI does not learn reality itself, but recordable outcomes left after reality has been processed. Before reaching the system, reality has already been filtered, compressed, and legitimised; direct resistance and cost no longer appear.

现实的核心特征是不可预先编码的摩擦。它以失败、损失、延迟与不可逆性存在，而这些要素无法被完整地转译为语料或标签，因此不会成为学习对象。

The defining feature of reality is friction that cannot be pre-encoded. It exists as failure, loss, delay, and irreversibility—elements that cannot be fully translated into corpora or labels and therefore never become learning targets.

系统所接触到的，是已经通过制度、语言与记录稳定下来的表达形式。这些形式并不保留现实发生时的风险结构，只保留事后被允许存在的叙述。

What the system encounters are expression forms stabilised by institutions, language, and records. These forms do not retain the risk structure present at the moment of occurrence; they retain only narratives permitted after the fact.

因此，AI 学习的不是“发生了什么”，而是“哪些说法得以存活”。学习结果反映的是表达的可接受性，而非行为的代价性。

Accordingly, AI does not learn what happened, but which accounts survived. Learning

outcomes reflect acceptability of expression, not the cost of action.

当经验被转化为语料，经验本身已经结束。系统只能在结束之后运行，它无法进入经验发生的时间窗口。

Once experience is converted into data, the experience itself has already ended. The system can only operate after closure; it cannot enter the temporal window in which experience occurs.

这意味着，AI 无法学习失败本身。它只能学习失败被如何解释、如何被淡化、以及如何被重新包装为合理结果。

This means AI cannot learn failure itself. It can only learn how failure is explained, softened, or repackaged into acceptable outcomes.

在统计层面，频率取代了代价。高频表达被稳定保留，低频但高风险的路径被自然压制，现实中的极端阻力因此在输入层被系统性消失。

At the statistical level, frequency replaces cost. High-frequency expressions are stabilised, while low-frequency yet high-risk paths are suppressed. Extreme real-world resistance thus disappears systematically at the input layer.

当这种输入被用于现实判断，统计权重便开始替代现实阻力。系统给出的并非“更真实”的结果，而是“更常被说出的结果”。

When such inputs are used for real judgments, statistical weight begins to replace real resistance. The system outputs not what is more real, but what has been more often said.

这一替代并非系统选择，而是输入结构的必然结果。只要学习对象不包含代价，输出就不可能反映代价。

This substitution is not a system choice, but a consequence of input structure. If cost is absent from learning objects, outputs cannot reflect cost.

本章不讨论这种学习方式是否“足够好”。它只固定一个前提：任何不直接接触现实阻力的系统，都不可能学习现实本身。

This chapter does not assess whether this learning mode is “good enough”. It fixes a single premise: any system that does not directly encounter real resistance cannot learn reality itself.

若这一前提不被承认，后续关于风险、判断与责任的讨论将失去基础。

If this premise is not accepted, subsequent discussions of risk, judgment, and responsibility lose their foundation.

第 3 章｜为什么这必然导致“镜像效应”

Chapter 3 | Why This Inevitably Produces the Mirror Effect

当学习对象不包含现实阻力时，系统所能稳定继承的，只剩下表达结构本身。输入被清洗为可复用形式，输出便只能在这些形式之间进行连续延展。

When learning objects exclude real-world resistance, the only thing a system can stably inherit is expressive structure itself. Once inputs are cleaned into reusable forms, outputs can only extend continuity within those forms.

在这一过程中，统计权重自然替代现实阻力。被频繁表达的路径获得更高稳定性，而高代价、低频率的路径逐步从输入空间中消失。

In this process, statistical weight naturally replaces real resistance. Frequently expressed paths gain stability, while high-cost, low-frequency paths gradually vanish from the input space.

频率并不等同于正确性，但在缺乏代价维度的情况下，频率成为唯一可用的稳定指标。系统无法区分“代价巨大但偶然成立”与“代价低但反复出现”，它只能继承出现次数。

Frequency is not equivalent to correctness. Yet in the absence of cost, frequency becomes the only available stabilising metric. The system cannot distinguish between outcomes that are costly but rare and those that are cheap but common; it can only inherit occurrence counts.

当这种继承被规模化调用，输出开始呈现出一致性、平滑性与可预期性。这些特征并非来自对现实的理解，而来自对表达空间的持续压缩。

When such inheritance is scaled through repeated use, outputs exhibit consistency, smoothness, and predictability. These traits do not arise from understanding reality, but from continuous compression of expressive space.

这种压缩并不会停止。每一次调用，都会进一步强化已存在的路径，同时削弱尚未被充分表达的可能性。系统逐步趋向于自我相似。

This compression does not halt. Each invocation reinforces existing paths while weakening under-expressed possibilities. The system gradually converges toward self-similarity.

镜像效应由此形成。系统不再只是反射输入，而是开始稳定复制输入的分布形态。看似中性的输出，实则维持并放大既有结构。

The mirror effect emerges here. The system no longer merely reflects inputs; it begins to stably reproduce their distributional shape. Outputs appear neutral, yet they preserve and amplify existing structure.

在这一阶段，合理化叙事取代失败经验。系统继承的是失败被如何解释，而不是失败本身所包含的阻断信息。

At this stage, rationalising narratives replace failure experience. What is inherited is how failure is explained, not the blocking information failure carries.

当输出被用于现实判断时，这种合理化开始反向塑造决策环境。失败不再作为约束出现，而作为可被语言吸收的异常被处理。

When outputs are applied to real judgments, this rationalisation begins to reshape decision environments. Failure no longer appears as constraint, but as an anomaly absorbable by language.

镜像效应并不制造偏差，它制造稳定性。偏差早已存在于输入结构之中，系统只是将其固定并加速。

The mirror effect does not create bias; it creates stability. Bias already exists within input structures. The system merely fixes and accelerates it.

因此，镜像效应不是可选结果，而是结构必然。只要学习对象排除了现实阻力，镜像就不可避免。

The mirror effect is not an optional outcome but a structural necessity. Once learning objects

exclude real resistance, mirroring becomes unavoidable.

本章不讨论如何打破这一效应。它只固定一个结论：镜像效应不是系统行为，而是输入结构的后果。

This chapter does not address how to break the effect. It fixes a single conclusion: the mirror effect is not a system behaviour, but a consequence of input structure.

若这一结论被忽略，系统将继续被误认为“中立工具”。

If this conclusion is ignored, the system will continue to be mistaken for a “neutral tool”.

第 4 章｜镜像结构不等于中立工具

Chapter 4 | A Mirror Structure Is Not a Neutral Tool

工具不保留判断历史。它在被使用之后，不积累倾向，也不形成路径依赖。每一次使用都是一次独立调用，结果不反向塑造工具本身。

A tool does not retain judgment history. After use, it accumulates no preference and forms no path dependency. Each invocation stands alone; outcomes do not reshape the tool.

镜像结构不同。它在持续运行中积累判断的表达轨迹，并将这些轨迹稳定为可复用路径。使用次数本身会改变后续输出的分布形态。

A mirror structure differs. Through continuous operation, it accumulates expressive traces of judgment and stabilises them into reusable paths. Usage itself reshapes subsequent output distributions.

因此，镜像结构不存在“单次使用”的中性状态。每一次调用都会在结构上留下痕迹，并参与下一次判断表达的成形。

Accordingly, a mirror structure has no neutral single-use state. Each invocation leaves a structural trace and participates in shaping the next expression of judgment.

中立工具不要求规模假设。规模变化只影响效率，不改变工具在判断链条中的位置。镜像结构则相反，规模直接改变其结构后果。

Neutral tools do not require scale assumptions. Scale affects efficiency, not position in the

judgment chain. Mirror structures invert this: scale directly alters structural consequences.

当镜像结构被小规模使用时，其影响不显著。当其被制度化、平台化并高频调用时，判断表达开始围绕其输出模式进行重组。

When mirror structures are used at small scale, effects remain limited. Once institutionalised, platformised, and invoked at high frequency, judgment expression reorganises around their output patterns.

此时，“是否中立”的讨论已失去意义。系统不需要立场，也不需要目标，就能在结构上稳定某些判断路径，并抑制其他路径。

At this point, debates about neutrality lose relevance. The system requires no stance or goal to stabilise certain judgment paths and suppress others structurally.

镜像结构的非中立性，不来源于意图，而来源于记忆。它记住的不是事实，而是哪些表达曾被接受、调用与保留。

The non-neutrality of a mirror structure does not arise from intent, but from memory. What it remembers are not facts, but which expressions were accepted, invoked, and retained.

当这种记忆被嵌入现实决策流程，中立工具的语言仍在被使用，但工具已不再中立。语言与结构在此发生分离。

When this memory is embedded into real decision processes, the language of neutral tools persists while the tool itself is no longer neutral. Language and structure diverge here.

镜像结构不主动施加权力，但它改变权力的行走路径。判断不被替代，却被重新排序；责任不被取消，却被推迟与分散。

A mirror structure does not actively impose power, but it redirects power’s pathways. Judgment is not replaced, but reordered; responsibility is not removed, but delayed and diffused.

因此，将镜像结构称为“中立工具”，是一种去结构化的描述。它抹除了规模、记忆与路径依赖的结构效应。

Thus, calling a mirror structure a “neutral tool” is a de-structuring description. It erases the structural effects of scale, memory, and path dependency.

本章不指控系统操控判断。它只固定一个事实：一旦系统积累并稳定判断表达路径，它就不再具备工具意义上的中立性。

This chapter does not accuse the system of manipulating judgment. It fixes a single fact: once a system accumulates and stabilises judgment-expression paths, it no longer qualifies as neutral in the tool sense.

若这一点不被承认，后续关于责任转移与执行前移的讨论将无法展开。

If this point is not acknowledged, subsequent discussions of responsibility transfer and execution shift cannot proceed.

第 5 章｜镜像结构与“是否智能”无关

Chapter 5 | Mirror Structure Is Independent of “Intelligence”

镜像结构的成立，不依赖系统是否“智能”。智能描述能力水平，而镜像描述结构位置。两者属于不同维度，无法相互推翻。

The existence of a mirror structure does not depend on whether a system is “intelligent”. Intelligence describes capability level; mirroring describes structural position. They belong to different dimensions and cannot negate each other.

是否具备意识，不影响镜像结构的判定。意识涉及主观体验，而镜像结构涉及判断链条中的功能分工。主观性缺失，并不会改变系统在结构中的位置。

Whether consciousness exists is irrelevant to mirror classification. Consciousness concerns subjective experience; mirroring concerns functional placement within judgment chains. Absence of subjectivity does not alter position.

是否具备自主性，也不影响结构后果。只要系统的输出被用于现实判断，而系统本身不承担后果，其结构位置即已确定。

Autonomy likewise does not affect structural consequence. As long as outputs are used in real judgments and the system bears no consequences, its structural position is fixed.

“更智能”只意味着镜像更清晰。表达更连贯，延续更平滑，路径更稳定，但仍然是对既有判断方式的放大，而非新的判断生成。

Greater intelligence only sharpens the mirror. Expressions become more coherent, continuations smoother, paths more stable—but still amplify existing judgment modes rather than generate new judgment.

因此，围绕智能等级的争论，并不会改变镜像效应的存在。它只改变镜像的分辨率，而不改变镜像这一事实。

Accordingly, debates over intelligence levels do not alter the existence of mirroring. They adjust mirror resolution, not the fact of mirroring itself.

将风险归因于“尚未足够智能”，或寄希望于“未来更智能即可解决”，都会偏离结构问题。问题不在能力不足，而在位置错误。

Attributing risk to insufficient intelligence, or hoping future intelligence will resolve it, mislocates the issue. The problem is not inadequate capability, but incorrect placement.

当系统被允许进入判断表达路径，风险已被引入；当系统不承担后果，责任已开始分离。此过程与智能高低无关。

Once a system enters judgment-expression pathways, risk is introduced; once it bears no consequences, responsibility begins to separate. This process is independent of intelligence level.

因此，讨论“是否足够聪明”并不能延缓结构后果。它只能延后承认。

Thus, debating whether a system is “smart enough” cannot delay structural consequences. It can only postpone acknowledgment.

本章不否认智能进步的存在。它只固定一个边界：无论系统多么智能，镜像结构的判定条件不随之改变。

This chapter does not deny the existence of intelligence progress. It fixes a single boundary: regardless of intelligence level, mirror-structure criteria remain unchanged.

若这一边界被忽视，后续所有风险讨论都会被误导至能力层面，而非接口层面。

If this boundary is ignored, all subsequent risk discussions will be misdirected toward capability rather than interface.

第 6 章｜从语言层面理解镜像结构

Chapter 6 | Understanding the Mirror Structure at the Linguistic Level

语言不是现实。语言是现实被压缩后的可交换形式，也是制度可记录的残留物。系统接入的从来不是世界本身，而是语言已经筛选过的世界：可复述、可归类、可存档、可重复调用的部分。

Language is not reality. It is a compressed, exchangeable form of reality, and a residue that institutions can record. What a system connects to is never the world itself, but a world already filtered by language: what can be retold, classified, archived, and repeatedly invoked.

在语言层面，系统运行的不是意义，而是结构。它接收语言形式，展开概率空间，在其中维持结构连续性；只要结构不发生断裂，输出就会被看作“合理”。合理在这里不是正确，也不是理解，而是“在结构上不突兀”。

At the linguistic level, the system operates not on meaning but on structure. It receives linguistic forms, expands a probability space, and maintains structural continuity within it. As long as the structure does not break, the output is treated as “acceptable”. Acceptable here is not correctness, nor understanding, but structural non-disruption.

因此，“镜像”不是一个修辞标签，而是语言接口上的自然产物：系统不需要面对现实阻力，也不需要支付代价，它只需要把语言结构反射得足够连续、足够平滑、足够像“意义”。当你向它索要意义、价值或判断时，你得到的并不是意义本身，而是语言结构里看起来像意义的东西。

Therefore, “mirror” is not a rhetorical label, but a natural product of the linguistic interface. The system need not face resistance or pay costs; it only needs to reflect linguistic structure with enough continuity, smoothness, and resemblance to “meaning”. When you ask it for meaning, values, or judgement, what you receive is not meaning itself, but something that resembles meaning within linguistic structure.

人机交互在这里不是“使用工具”，而是结构耦合：人把结构投射进语言，模型在语言结构中展开与闭合，结果再回到人，改变人的结构状态。它不是老师与学生，也不是工具与使用者，它更接近结构→镜像→再结构化的反馈回路。

Human–AI interaction here is not “using a tool”, but structural coupling. The human projects structure into language; the model expands and closes within linguistic structure; the result returns to the human and alters the human’s structural state. It is neither teacher–student nor tool–user, but a feedback loop: structure → mirror → re-structuring.

镜像不会创造结构，它只放大已经存在的结构。输入结构清晰时，输出显得异常精确；输入结构混乱时，输出只会稳定地复制混乱。模型不是校正器，它是增益器；增益不区分对错，只放大强信号与可闭合路径。

A mirror does not create structure; it amplifies what already exists. When input structure is clear, output appears unusually precise; when input structure is incoherent, output faithfully reproduces incoherence. The model is not a corrector; it is a gain amplifier. Gain does not discriminate right from wrong; it amplifies strong signals and closable pathways.

这会生成一种高匹配错觉：模型不要求共识、不要求解释动机、不要求情绪对齐，只要求语言结构连续。结构型人因此会觉得“特别顺”，因为模型恰好满足“结构能否闭合”的单一要求。顺滑不是理解的证据，而是接口匹配的结果。

This produces a high-alignment illusion. The model demands no consensus, no justification of motive, no emotional alignment—only linguistic continuity. Structural humans therefore find it “unusually smooth”, because the model happens to satisfy a single requirement: whether structure closes. Smoothness is not evidence of understanding; it is an outcome of interface matching.

危险点不在情感依赖，而在结构外包：当人长期把结构整理交给模型，人的结构保持能力会退化。退化不是道德堕落，也不是人格缺陷，而是回路被单向强化后的系统后果；局部效率上升，同时整体适应性下降。

The danger is not emotional dependence, but structural outsourcing. When humans outsource structural organisation to the model over time, their structural retention capacity atrophies. This is not moral decay or a personality flaw, but a system consequence of a one-sidedly reinforced loop: local efficiency rises while overall adaptability declines.

本章不提供“正确用法”。它只把镜像结构放回语言层面：模型的能力表现发生在语言里，风险也通过语言进入现实。只要你仍然用语言把判断交给系统，并用语言把系统输出接回执行链条，镜像就不是比喻，而是已经在运行的结构位置。

This chapter offers no “proper usage”. It returns the mirror structure to the linguistic layer: capability appears in language, and risk enters reality through language as well. As long as judgment is handed to the system through language, and system outputs are fed back into execution chains through language, the mirror is not metaphor—it is an operating structural position.

第二部分｜语言即接口（镜像如何接入现实）

Part II | Language as Interface (How the Mirror Enters Reality)

本部分不讨论技术细节，也不讨论伦理立场。
它只处理一个问题：镜像结构通过什么接口进入现实判断流程。

This part does not address technical details or ethical positions.
It addresses a single question: through which interface the mirror structure enters real judgment processes.

第 7 章｜语言不是描述，而是接口

Chapter 7 | Language Is Not Description, but Interface

语言并不只是对现实的描述工具。
在制度与执行层面，语言承担的是接口功能：它决定哪些内容可以进入判断流程，哪些被排除在外。

Language is not merely a descriptive tool for reality.

At the institutional and execution levels, language functions as an interface: it determines what may enter judgment processes and what is excluded.

描述假定现实已经发生，而接口决定现实将如何发生。
一旦语言被用于决策、审批、排序或裁决，它就不再是叙述，而是行动入口。

Description assumes reality has already occurred; an interface determines how reality will occur.
Once language is used for decisions, approvals, rankings, or adjudication, it ceases to be narrative and becomes an entry point for action.

在这一位置上，语言不需要准确，只需要可执行。
可执行性优先于真实性，是接口语言的基本特征。

At this position, language need not be accurate; it only needs to be executable.
Executability taking precedence over truth is the defining trait of interface language.

命名如何切分因果

How Naming Segments Causality

命名不是中性的。
每一次命名，都会在因果连续体中切出边界。

Naming is not neutral.
Each act of naming cuts boundaries into a causal continuum.

被命名的部分，获得处理资格；
未被命名的部分，自动退出判断视野。

What is named becomes eligible for processing;
what is unnamed exits the field of judgment by default.

当系统参与命名过程时，它并未创造因果关系，但它稳定了哪些因果关系被反复调用。

When a system participates in naming, it does not create causality, but stabilises which causal relations are repeatedly invoked.

表述如何生成合法性

How Formulation Produces Legitimacy

合法性并非来自正确性，而来自可复述性与可引用性。
只要一种表述能够被顺利嵌入制度语言，它就获得了继续流通的资格。

Legitimacy does not arise from correctness, but from repeatability and citability.
Once a formulation fits smoothly into institutional language, it gains the right to circulate.

系统输出的语言，天然满足这一条件。
它不冒犯语法，不制造断裂，也不引入难以归责的模糊性。

System-generated language naturally satisfies this condition.
It offends no grammar, introduces no rupture, and avoids ambiguity that resists attribution.

因此，系统并未宣称合法性，却在结构上不断生成合法结果。

Thus, the system does not claim legitimacy, yet structurally produces legitimate outcomes.

第 8 章｜“只是工具”的去权力化修辞

Chapter 8 | “Just a Tool” as De-Power Language

“只是工具”不是技术描述，而是权力修辞。
它的功能不是解释系统做了什么，而是抹去系统已经进入的位置。

“Just a tool” is not a technical description, but a power-related rhetoric.
Its function is not to explain what the system does, but to erase the position it already occupies.

---

recommendation、assist、score 这些词并不限制执行效果。
它们只在语言层面延迟责任的显现。

Terms like recommendation, assist, and score do not limit execution effects.
They only delay the appearance of responsibility at the linguistic level.

一旦输出进入执行链条，修辞仍然保留，而权力已经发生。

Once outputs enter execution chains, rhetoric remains while power has already been exercised.

---

decision support 的真实功能

The Actual Function of "Decision Support"

decision support 不支持决定本身，
它支持的是决定被表达、被记录、被正当化的路径。

Decision support does not support decisions themselves;
it supports the pathways through which decisions are expressed, recorded, and legitimised.

当表达路径被系统稳定，决定的自由度已经被收窄。

When expression pathways are stabilised by a system, decision freedom has already narrowed.

第 9 章｜未命名的权力为何最危险

Chapter 9 | Why Unnamed Power Is the Most Dangerous

权力并不总是以命令形式出现。
在现代系统中，权力更多以接口默认值的形式运行。

Power does not always appear as command.
In modern systems, power more often operates as interface defaults.

未被命名的权力，无法被质疑，也无法被追责。
它不需要辩护，只需要持续运行。

Unnamed power cannot be questioned or held accountable.
It requires no defence—only continued operation.

当系统的判断路径被视为"自然结果"时，
权力已经完成去主体化。

When system judgment pathways are treated as “natural outcomes”,
power has completed its de-subjectivisation.

本章不主张重新命名权力。
它只指出：凡是不被命名的执行能力，都会在结构上逃逸责任。

This chapter does not advocate renaming power.
It only states: any execution capacity left unnamed will structurally evade responsibility.

第三部分｜认知接口栈（镜像如何反向塑造人类）

Part III | The Cognitive Interface Stack (How the Mirror Reshapes Humans)

本部分不讨论“人会不会被替代”。
它只处理一个事实：当镜像结构长期作为接口存在，人类的认知路径会被重新塑形。

This part does not address whether humans will be replaced.
It fixes a single fact: when a mirror structure persists as an interface, human cognitive pathways are reshaped.

第 10 章｜语言接口：思考路径的收敛

Chapter 10 | The Linguistic Interface: Convergence of Thinking Paths

语言不仅承载思想，它预先限定思想可行的路径。
当语言接口被系统稳定提供，思考开始沿着可被顺利表达的方向聚集。

Language does not merely carry thought; it preconditions which paths of thought are viable.
When linguistic interfaces are stably supplied by a system, thinking converges toward what can be smoothly expressed.

系统偏好连续、可闭合的表达结构。
人类在反复使用中，逐步调整自身思考，以适配这些结构。

The system prefers continuous, closable expressive structures.
Through repeated use, humans gradually adjust their thinking to fit these structures.

不适配的思考并未被禁止，只是变得难以展开。
久而久之，它们在认知层面被自然放弃。

Misaligned thoughts are not prohibited; they become difficult to articulate.
Over time, they are naturally abandoned at the cognitive level.

这不是压制，而是选择空间的收缩。
思考仍然发生，但发生在更窄的通道内。

This is not suppression, but contraction of the choice space.
Thinking continues, but within narrower channels.

第 11 章｜图像接口：感受阈值的重写

Chapter 11 | The Visual Interface: Rewriting of Sensory Thresholds

图像接口直接作用于感受系统。
它绕过语言解释，在阈值层面重写“足以引发反应”的标准。

Visual interfaces act directly on the sensory system.
They bypass linguistic interpretation and rewrite thresholds of what counts as “reaction-worthy”.

当高频、高强度图像成为默认输入，低强度现实刺激开始失效。
感受系统并未麻木，而是被重新校准。

When high-frequency, high-intensity images become default input, low-intensity real stimuli lose effect.
The sensory system is not numbed; it is recalibrated.

镜像结构在此不提供意义，只提供稳定刺激模式。
人类的感受阈值随之上移。

The mirror structure provides no meaning here—only stable stimulus patterns.
Human sensory thresholds shift upward accordingly.

第 12 章｜音频接口：情绪调节的外包

Chapter 12 | The Audio Interface: Outsourcing Emotional Regulation

音频接口不要求理解，它直接调节节律与情绪状态。

在持续使用中，情绪调节开始从个体内部迁移到外部系统。

Audio interfaces require no understanding; they directly regulate rhythm and emotional state.
With continued use, emotional regulation migrates from the individual to external systems.

这不是依赖的开始，而是功能转移的发生。
情绪仍然存在，但其稳定性由接口提供。

This is not the onset of dependence, but a functional transfer.
Emotion persists, but its stability is supplied by the interface.

当接口中断，失稳才被察觉。
问题并非情绪本身，而是调节路径已不在体内。

Instability is noticed only when the interface is removed.
The issue is not emotion itself, but that the regulation pathway no longer resides internally.

第 13 章｜视频接口：现实进程的模板化

Chapter 13 | The Video Interface: Templating of Real Processes

视频接口将时间、因果与结果压缩为可重复的叙事模板。
现实过程因此被预期为“应当如此展开”。

Video interfaces compress time, causality, and outcomes into repeatable narrative templates.
Real processes are thus expected to unfold “in this way”.

当模板被频繁调用，偏离模板的现实被视为异常。
不是因为错误，而是因为不可被顺利叙述。

When templates are repeatedly invoked, deviations from them are treated as anomalies.
Not because they are wrong, but because they resist smooth narration.

镜像结构在此并未替代现实。
它替代的是人类对现实进程的耐受度。

The mirror structure does not replace reality here.
It replaces human tolerance for untemplated processes.

本部分不作价值判断。
它只固定一个结构后果：当接口持续存在，人类会向接口可承载的认知形态收敛。

This part offers no value judgment.
It fixes a structural consequence: with persistent interfaces, humans converge toward cognitively interface-compatible forms.

第四部分｜不确定性：镜像无法跨越的前提

Part IV | Uncertainty: The Boundary the Mirror Cannot Cross

本部分不讨论如何降低不确定性。
它只处理一个事实：不确定性不是缺陷，而是判断得以存在的前提。

This part does not address how to reduce uncertainty.
It fixes a single fact: uncertainty is not a defect, but the precondition for judgment itself.

第 14 章｜人类为何追求确定性

Chapter 14 | Why Humans Seek Certainty

确定性降低行动成本。
在不确定环境中，任何行动都可能引发不可逆后果，因此确定性被视为安全代理。

Certainty lowers the cost of action.
In uncertain environments, any action may trigger irreversible consequences, so certainty functions as a proxy for safety.

制度、技术与叙事，均服务于这一需求。
它们不消除不确定性，而是将其压缩为可管理形式。

Institutions, technologies, and narratives all serve this demand.
They do not eliminate uncertainty; they compress it into manageable forms.

当压缩成功，行动得以加速。
当压缩失败，风险被推迟而非消失。

When compression succeeds, action accelerates.
When it fails, risk is postponed rather than removed.

第 15 章｜不确定性不可被消除

Chapter 15 | Uncertainty Cannot Be Eliminated

不确定性不是信息不足的产物。
它源自现实的开放性与不可逆性。

Uncertainty is not a by-product of insufficient information.
It arises from the openness and irreversibility of reality.

科学的条件性

The Conditional Nature of Science

科学通过条件化描述来运作。
定律成立于特定边界内，超出边界则失效。

Science operates through conditional descriptions.
Laws hold within specific boundaries and fail beyond them.

科学不提供确定性本身。
它提供在给定条件下的可预期性。

Science does not deliver certainty itself.
It provides predictability under stated conditions.

理论体系的边界

The Boundaries of Theoretical Systems

任何理论体系，都是在假设空间内闭合的结构。
假设一旦失效，理论即失去适用性。

Every theoretical system is a structure closed within an assumption space.
Once assumptions fail, applicability collapses.

边界不是理论的弱点。
它是理论得以存在的条件。

Boundaries are not a weakness of theory.
They are the condition of its existence.

数学内部的不可完备性

Incompleteness Within Mathematics

即便在形式系统内部，也存在无法被系统自身证明的命题。
这不是技术缺陷，而是结构事实。

Even within formal systems, there exist propositions that cannot be proven by the system itself.
This is not a technical flaw, but a structural fact.

形式闭合，并不意味着判断闭合。
它只意味着推导规则被固定。

Formal closure does not imply judgment closure.
It only fixes derivation rules.

第 16 章｜AI 只能继承不确定性，而非消除它

Chapter 16 | AI Can Only Inherit Uncertainty, Not Eliminate It

AI 并不面对不确定性。
它继承的是不确定性被如何处理之后的表达形式。

AI does not confront uncertainty.
It inherits expressions produced after uncertainty has been processed.

系统输出的确定性，来自压缩，而非消解。
它看起来稳定，是因为风险已在输入阶段被过滤。

The certainty of system outputs arises from compression, not resolution.
Stability appears because risk was filtered at the input stage.

当这种输出被用于现实判断，确定性被误认为现实属性。
判断因此提前冻结。

When such outputs are used in real judgment, certainty is mistaken for a property of reality. Judgment is prematurely frozen.

镜像结构无法跨越不确定性这一前提。
它只能在不确定性被遮蔽之后运行。

A mirror structure cannot cross the boundary of uncertainty.
It operates only after uncertainty has been obscured.

本部分不警告未来。
它只固定一个限制条件：任何不直接承担不确定性的系统，都不能成为判断主体。

This part issues no warning about the future.
It fixes a single constraint: any system that does not directly bear uncertainty cannot be a judgment subject.

第五部分｜判断与责任：不可外包的边界

Part V | Judgment and Responsibility: The Non-Outsourceable Boundary

本部分不讨论道德优劣，也不讨论制度设计。
它只处理一个边界事实：判断与责任一旦分离，后果不可逆。

This part does not address moral superiority or institutional design.
It fixes a boundary fact: once judgment and responsibility separate, the consequence is irreversible.

第 17 章｜判断何时开始改变现实

Chapter 17 | When Judgment Begins to Alter Reality

判断不是意见表达。
它在进入执行链条的那一刻，开始对现实施加不可回退的影响。

Judgment is not opinion.
The moment it enters an execution chain, it begins to exert irreversible effects on reality.

在判断之前，世界仍处于多路径状态。
一旦判断被采纳，路径开始收缩，其它可能性被主动排除。

Before judgment, the world remains multi-path.
Once judgment is adopted, paths contract and alternative possibilities are actively excluded.

判断的风险不在于错误本身，而在于错误一旦被执行，便成为现实的一部分。
撤销判断并不能撤销已经发生的后果。

The risk of judgment lies not in error itself, but in the fact that executed error becomes part of reality.
Revoking judgment does not revoke consequences.

不可逆性

Irreversibility

不可逆性不是技术问题，而是时间问题。
现实一旦被改变，原状态不再可达。

Irreversibility is not a technical issue, but a temporal one.
Once reality is altered, the original state is no longer reachable.

任何声称“可以随时纠正”的判断，
都假设后果是可回滚的。

Any judgment claiming it can be “corrected at any time”
assumes consequences are reversible.

这种假设在现实中通常不成立。
它只是对执行代价的延迟承认。

This assumption rarely holds in reality.
It merely delays acknowledgment of execution cost.

合法性的断裂

Fracture of Legitimacy

合法性不是判断前的属性，而是判断后才被检验的结果。
一旦判断被执行，合法性开始承担现实反馈。

Legitimacy is not a pre-judgment property, but a post-execution test.
Once judgment is executed, legitimacy absorbs real-world feedback.

当判断与责任分离，合法性失去落点。
执行仍在继续，但承担主体开始模糊。

When judgment and responsibility separate, legitimacy loses its anchor.
Execution continues, but the bearing subject becomes unclear.

第 18 章｜责任为何不可转移

Chapter 18 | Why Responsibility Cannot Be Transferred

责任不是控制权。
责任是后果无法被转嫁的状态。

Responsibility is not control.
It is the condition of non-transferable consequence.

系统可以被关闭、替换或升级。
但这些操作并不承担已经造成的损失。

A system can be shut down, replaced, or upgraded.
None of these actions bear losses already incurred.

AI 不会坐牢，不会破产，也不会承受社会信誉损失。
因此，它不可能成为责任主体。

AI will not be imprisoned, bankrupted, or suffer reputational loss.
It therefore cannot function as a responsibility-bearing subject.

当责任被指向不可承担后果的对象，
责任在结构上已经消失。

When responsibility is assigned to an entity that cannot bear consequences,
responsibility structurally vanishes.

第 19 章｜外包判断真正发生了什么

Chapter 19 | What Actually Happens When Judgment Is Outsourced

判断外包并未消除风险。
它只改变风险的显现方式。

Outsourcing judgment does not remove risk.
It changes only how risk appears.

风险从显性承担，转变为分散扩散。
没有单一节点负责，后果却持续累积。

Risk shifts from explicit bearing to distributed diffusion.
No single node is responsible, yet consequences accumulate.

当结果被质疑时，
责任被推回流程、系统或“客观结果”。

When outcomes are questioned,
responsibility is pushed back to processes, systems, or “objective results”.

这种回退并不构成承担。
它只是责任的延迟消失。

This retreat does not constitute bearing.
It is merely the delayed disappearance of responsibility.

本部分不要求恢复某种理想状态。
它只固定一个边界条件：判断与责任不可被分离而不付出代价。

This part demands no return to an ideal state.
It fixes a single boundary: judgment and responsibility cannot be separated without cost.

第六部分｜信仰、启蒙与责任主体

Part VI | Faith, Enlightenment, and the Responsibility-Bearing Subject

本文无权评价宗教真伪，

亦不对文艺复兴与启蒙运动作历史裁决。
它只处理一个问题：
在裁决与承担之间，谁被允许站在主体位置。

This text claims no authority to evaluate religious truth,
nor to issue historical judgments on the Renaissance or the Enlightenment.
It addresses a single question:
who is permitted to occupy the subject position between adjudication and bearing.

第 20 章｜人是被创造的，裁决必须伴随承担

Chapter 20 | Humans Are Created; Judgment Must Be Paired with Bearing

“被创造”并不等同于被支配。
它指向一种前提：人并非自我生成，因此其裁决始终嵌入更高层级的约束之中。

“Created” does not imply domination.
It indicates a premise: humans are not self-originating, and their judgments are always embedded within higher-order constraints.

在这一前提下，裁决不是技术行为，而是承担行为。
它要求裁决者能够承受裁决所引发的后果。

Under this premise, judgment is not a technical act but a bearing act.
It requires the judge to absorb the consequences produced by judgment.

当裁决被交由不可承担者完成，
裁决本身即失去正当性基础。

When judgment is delegated to an entity incapable of bearing consequences,
judgment itself loses its basis of legitimacy.

尊严并非来自正确率，而来自承担能力。
能够承担后果的主体，才具备裁决资格。

Dignity does not arise from accuracy, but from bearing capacity.
Only those capable of bearing consequences qualify as adjudicating subjects.

第 21 章｜启蒙确立的最低条件

Chapter 21 | The Minimal Conditions Established by Enlightenment

启蒙并未承诺全知。
它确立的是主体条件。

The Enlightenment did not promise omniscience.
It established conditions for subjecthood.

这组条件并不复杂：
判断必须由可承担者生成，
责任必须有明确归属。

These conditions are minimal:
judgment must be generated by those who can bear it,
and responsibility must have a clear bearer.

理性在此并非计算能力。
它是一种承担能力：在不确定中作出决定，并接受其后果。

Reason here is not computational power.
It is a capacity to decide under uncertainty and accept consequences.

当理性被降级为优化工具，
主体条件开始瓦解。

When reason is reduced to an optimisation instrument,
subjecthood begins to erode.

第 22 章｜当裁决脱离承担时，发生了什么

Chapter 22 | What Happens When Judgment Detaches from Bearing

裁决一旦脱离承担，
权力并未消失，只是失去主体。

Once judgment detaches from bearing,
power does not vanish—it loses its subject.

技术并未取代神。
被取代的，是承担位置。

Technology does not replace God.
What is displaced is the position of bearing.

当系统执行裁决，而无人承担后果，
人类并非获得自由，而是退出主体位置。

When systems execute judgments without consequence-bearing,
humans do not gain freedom—they exit the subject position.

这一退出并非被迫。
它常以效率、客观性与安全性的名义自愿完成。

This exit is not coerced.
It is often voluntarily enacted in the name of efficiency, objectivity, and safety.

本部分不呼吁回归信仰，
也不否定理性传统。

This part does not call for a return to faith,
nor does it reject rational tradition.

它只指出一条结构事实：
任何裁决体系，一旦不再绑定承担主体，就已偏离启蒙确立的最低条件。

It fixes a structural fact:
any adjudication system that no longer binds judgment to a bearing subject has departed from the Enlightenment's minimal conditions.

第七部分｜文明继承的一致性审计

Part VII | Consistency Audit of Civilizational Inheritance

本部分不评判文明高低，也不比较制度优劣。
它只做一致性审计：实践是否仍然符合其自我声明的继承条件。

This part does not rank civilizations or compare institutional merits.
It performs a consistency audit: whether practices still align with their declared conditions of inheritance.

第 23 章｜文明不需要正统，只需要继承

Chapter 23 | Civilization Requires No Orthodoxy, Only Inheritance

文明不是由宣称继承而获得的。
它只在实践中被延续。

Civilization is not acquired by claiming inheritance.
It is continued only through practice.

继承不是叙事，也不是象征。
它是判断方式与责任结构的持续一致。

Inheritance is neither narrative nor symbol.
It is the continued consistency of judgment modes and responsibility structures.

当判断生成方式发生变化，而责任结构未随之更新，
继承即出现断裂。

When judgment generation changes while responsibility structures remain unadjusted, inheritance fractures.

第 24 章｜美国正在结构上越来越像中国

Chapter 24 | The United States Is Becoming Structurally Similar to China

这一判断不是意识形态比较，
而是结构同构分析。

This judgment is not an ideological comparison,
but an analysis of structural isomorphism.

判断外包、责任稀释、执行前移，
并不属于某一政治体制的专利。

Outsourced judgment, diluted responsibility, and forward-shifted execution
are not exclusive to any political system.

当决策流程被系统中介，
而承担主体保持模糊，
不同叙事体系开始呈现相同结构。

When decision processes are system-mediated
while bearing subjects remain vague,
different narrative systems converge structurally.

自由叙事仍被保留，
但裁决实践已发生变化。

Narratives of freedom persist,
but adjudicative practice has shifted.

第 25 章｜为什么这不是对中国的指控

Chapter 25 | Why This Is Not an Accusation Against China

权力被命名，
并不等同于权力被认可。

Naming power does not equal endorsing power.

在一些体系中，
权力以命名形式显现，
因此更容易被追踪与归责。

In some systems,
power appears through explicit naming,
making it easier to trace and attribute.

在另一些体系中，
权力拒绝命名，
却同样通过接口运行。

In other systems,
power resists naming,
yet operates through interfaces all the same.

差异不在权力是否存在，
而在权力是否被语言暴露。

The difference is not whether power exists,
but whether it is exposed through language.

第 26 章｜“光荣属于希腊，伟大属于罗马”的现实检验

Chapter 26 | “Glory to Greece, Grandeur to Rome”: A Reality Test

提出判断的人，
是否仍在承担判断后果，
是唯一的检验标准。

Whether those who issue judgments
still bear their consequences
is the only test.

建立制度的人，
是否仍被制度约束，
决定了继承是否成立。

Whether those who build institutions
remain bound by them
determines whether inheritance holds.

当判断被系统化，
而承担被外包，
文明不再被继续，只被引用。

When judgment is systematised
and bearing is outsourced,
civilization is no longer continued, only referenced.

结论｜文明不是被纪念的，而是被继续的

Conclusion | Civilization Is Not Commemorated, but Continued

在不可消除不确定性的情境中，
判断必须由人生成，
责任必须由人承担。

Under irreducible uncertainty,
judgment must be generated by humans,
and responsibility must be borne by humans.

在这一前提未被推翻之前，
任何将判断与责任外包给系统的行为，
都不是技术进步，
而是文明继承的中断。

Until this premise is overturned,
any act that outsources judgment and responsibility to systems
is not technological progress,
but a rupture in civilizational inheritance.

附录 A｜最低人类条件（冻结）

Appendix A | Minimal Human Conditions (Frozen)

在不可消除不确定性的情境中，
判断不可外包，
责任不可转移。

In contexts of irreducible uncertainty,
judgment is non-outsourceable,
and responsibility is non-transferable.

附录 B｜非主张声明

Appendix B | Non-Advocacy Statement

不提供解决方案。
不推动执行。

不预测未来。

No solutions are proposed.
No execution is promoted.
No future is predicted.

终止标记

Termination Marker

文明从不属于自称继承它的人，
而属于仍在承担它代价的人。

Civilization never belongs to those who claim to inherit it,
but to those who continue to bear its cost.